

Lecture 2: Supervised Learning

Instructor: Weinan E

Scribe: Guanhua Huang, Zhong Zheng

1 Introduction to Machine Learning

- **Statistics:** less focus on algorithm, more model design & hypothesis testing)
- **Machine Learning:** less focus on hypothesis testing, more on algorithm
 - **supervised learning:**
 - * **Data:** $\{x_j, y_j\}_{j=1}^n$
 - * **Model:** $y_j = f^*(x_j) + \epsilon_j$ where y_j is label, ϵ_j is noise.
 - * **Objective:** *approximate* f^*
 - **unsupervised learning:**
 - * **Data:** $\{x_j\}_{j=1}^n$
 - * **Model:** $y_j = f^*(x_j) + \epsilon_j$ where y_i is label
 - * **Objective:** find out the rule
 - **reinforcement learning:** policy function: from state space to action space
 - * **Objective:** optimal decision

2 Supervised Learning

- **Regression:** $f^* : D \subset R^d \rightarrow R$ where f^* is continuous.
- **Classification:** $f^* : D \rightarrow G(\text{finite set}), G = \{-1, 1\}$
- **Framework**
 - **Hypothesis Space \mathcal{H}_m**
 - * e.g. $\mathcal{H}_m = \{w_0 + w^\top x, w_0 \in R, w \in R^d\}$
 - * e.g. $\mathcal{H}_m = \{\sum_{j=1}^m \alpha_j \phi_j(x)\}$ where $\{\phi_j(x)\}_{j=1}^m$ is fixed set of function. For example, we can set $\phi_j(x) = \cos(k_j x)$
 - * e.g. $\mathcal{H}_m = \{\sum_{j=1}^m a_j \sigma(b_j^\top x + c_j)\}$ For example, we can set $\sigma(x) = \max\{0, x\}$ (ReLU, an activation function for Neural Network)
 - * In examples, m represents (scale of) degree of freedom (You can search **VC dimension** if you want know more about it)

- **Objective Function:** (Loss Function): loss function here is an example of square loss

$$\hat{R}_n(\theta) = \frac{1}{n} \sum_{j=1}^n (\hat{y}_j - f(x_j, \theta))^2 + \lambda \|\theta\|$$

Where θ is parameter, $\|\theta\|$ is norm, $\lambda \|\theta\|$ is regularization term.

- **Optimization Method**

- * Gradient Decent
- * SGD, Adam
- * BFGS .etc

3 Examples for Supervised Learning

3.1 Linear Model

$$\mathcal{H}_m = \{w_0 + w^\top x, w_0 \in R, w \in R^d\}$$

write $w_0 + w^\top x$ as $w^\top x$

$$\hat{R}_n(\theta) = \frac{1}{n} \sum_{j=1}^n (w^\top x_j - y_j)^2$$

$$\nabla_{\theta} \hat{R}_n(\theta) = \sum (w^\top x_j - y_j) x_j = 0$$

$$X = (x_1, \dots, x_n)$$

$$\hat{w} = (XX^\top)^{-1} Xy$$

Regularization **Ridge Regression**

$$\hat{R}_n(\theta) = \frac{1}{n} \sum_{j=1}^n (w^\top x_j - y_j)^2 + \lambda \|w\|^2$$

$$\hat{w} = (XX^\top + \lambda I)^{-1} Xy$$

$$\lim_{\lambda \rightarrow 0} (XX^\top + \lambda I)^{-1} = (XX^\top)^{-1} \text{ (Generalized inverse matrix)}$$

Use another regularization term

$$\hat{R}_n(\theta) = \frac{1}{n} \sum_{j=1}^n (w^\top x_j - y_j)^2 + \lambda N(w)$$

$$N(w) = \text{number of non-zero component of } w$$

Using $\|w\|_1$, Model become **Lasso Regression**

$$\hat{R}_n(\theta) = \frac{1}{n} \sum_{j=1}^n (w^\top x_j - y_j)^2 + \lambda \|w\|_1$$

Dimension Reduction of Feature Space with LASSO




Figure source: <https://towardsdatascience.com>

3.2 Kernel Method

- **kernel function** $k(x, y), x, y \in R^d$, e.g.

$$k(x, y), x, y \in R^d$$

- **Definition:**

- k is symmetric $k(x, y) = k(y, x)$
- $\forall \{x_j\}_{j=1}^n K = (k(x_i, x_j))_{n \times n} \geq 0$ K is **SPD**(symmetric positive definite)

- **Kernel Space:**

$$\mathcal{H}_m = \left\{ \sum_{j=1}^n \alpha_j k(x_j, x) \right\} (m = n)$$

- **Feature-based method:** $\{\phi_j(x)\}_{j=1}^m$ is a set of features

$$\mathcal{H}_m = \left\{ \sum_{i=1}^m \alpha_i \phi_i(x) \right\}$$

3.3 Neural Network

$$\mathcal{H}_m = \left\{ \sum_{j=1}^m a_j \sigma(b_j^\top x) \right\}$$




Figure source: <https://www.datasciencecentral.com/profiles/blogs>
Why NN better? (No complete theory in mathematics)

Compared with generalized linear model(GLM)

$$\mathcal{H}_m = \left\{ \sum_{i=1}^m \alpha_j \phi_j(x) \right\}$$

We have a "**Theorem**", For GLM

$$f : R^d \rightarrow R, d \gg 1 \inf_{f_m \in \mathcal{H}_m} \|f_m - f\| \geq cm^{-\frac{1}{d}}$$

For NN

$$f : R^d \rightarrow R, d \gg 1 \inf_{f_m \in \mathcal{H}_m} \|f_m - f\| \leq cm^{-\frac{1}{2}}$$

So let $error = 0.1$. For NN

$$m \sim 10^2$$

For GLM

$$m \sim 10^d$$

3.4 Optimization Algorithm

Gradient Descent:

$$\min_{\theta} F(\theta) \theta_{k+1} = \theta_k - \eta_k \nabla F(\theta_k)$$

For example

$$\nabla F(\theta) = \frac{1}{2} \sum_{j=1}^n \nabla (f(\theta, x_j) - y_j)^2 = \sum (f(\theta, x_j) - y_j) \nabla_{\theta} f$$

Back Propagation : Just use **Chain Rule** In practice, we use Stochastic Gradient Descent (SGD) Methods because of the large scale of data.

3.5 Classification

$$y = \{-1, 1\}$$

$$y = H(f(x))H(z) = \begin{cases} 1, z > 0 \\ 0, z \leq 0 \end{cases}$$

For continuity, we can let $H(z) = \frac{1}{1+e^{-z}}$ (sigmoid)

Logistic Regression: set $f = w^{\top} x$ and $y = \frac{1}{1+e^{-w^{\top} x}}$ Above are binary classifications. For **multi-class classification** (K classes), we use softmax

$$q_j(x) = \frac{e^{f_j(x)}}{\sum_{k=1}^K e^{f_k(x)}} \sum_{j=1}^K q_j(x) = 1$$

Where $q_j(x)$ can be regard as the probability of $x \in class_j$